


Table of Contents

- **EloqKV**
 - Introduction
 - Architecture
 - Core Features
 - Benchmark Report
 - Use Case



EloqKV: Redis-Compatible Key-Value Store with Memory-Like Performance and SSD Economics

Introduction

Artificial intelligence is rapidly becoming one of the largest consumers of computing infrastructure ever deployed. Training and inference systems now rely on massive GPU clusters operating at a scale that did not exist just a few years ago. While much of the attention has focused on compute, these systems are also driving unprecedented demand for high-performance memory, placing new pressure on the global DRAM supply chain.

The result is a dramatic shift in the economics of memory. As AI infrastructure expands, large portions of global DRAM production are being allocated to GPU-centric systems, tightening supply across the broader technology market. In 2025, average DRAM selling prices increased by more than 150 percent in some segments, and analysts expect continued pressure into 2026. In many cases, organizations are finding it difficult to acquire additional memory capacity even when they are willing to absorb the higher costs.

As a result, engineering teams are beginning to reconsider architectures that depend heavily on large in-memory data sets. Systems designed around abundant and inexpensive DRAM are now encountering new economic constraints. When memory was plentiful, this design approach was rarely questioned. Today, those assumptions are beginning to change.

The Limits of DRAM-First Architectures

Many modern applications achieve extremely low latency by keeping active datasets entirely in memory. Technologies such as Redis popularized this model by enabling fast access to data structures stored in DRAM.

Because Redis was designed as an in-memory system, practical dataset scale is closely tied to the amount of memory available to store the active working set. As datasets grow, additional DRAM must be provisioned to maintain performance. In an environment where memory is becoming both more expensive and harder to obtain, this dependency is beginning to introduce new economic limits on system scale.

These changes raise a fundamental question: how can high-performance data systems maintain low latency without requiring the entire dataset to reside in memory?

Introducing EloqKV

EloqKV was designed to address this challenge. Rather than assuming that high-performance data systems must keep the entire active dataset in DRAM, EloqKV uses an SSD-first architecture that leverages the performance of modern NVMe storage.

By combining Redis compatibility with a storage architecture optimized for solid-state media, EloqKV enables applications to manage significantly larger datasets while reducing dependence on expensive and increasingly constrained DRAM resources.

EloqKV can also utilize DRAM to cache extremely hot data, ensuring ultra-low latency while allowing the majority of the dataset to reside on high-performance SSDs.


All in all, by replacing expensive DRAM with NVMe SSDs for the majority of the data footprint, EloqKV removes the hard capacity ceiling of DRAM-based systems while delivering up to **10x lower total cost of ownership (TCO)** for large-scale workloads.

Architecture

EloqKV is a Redis-compatible distributed key-value store powered by Data Substrate. Its architecture includes a frontend compute engine that is fully compatible with the Redis protocol. Within Data Substrate, the TxService is responsible for caching hot data and managing transaction processing, while the LogService handles data persistence.

LogService replicas are distributed across multiple availability zones (AZs) to ensure resilience against AZ-level failures. The underlying storage layer supports pluggable key-value (KV) stores, such as EloqStore, RocksCloud, AWS DynamoDB, Google Bigtable, and Cassandra. These storage systems hold cold data for cache misses and provide high availability for baseline data.

Our preferred storage engine, EloqStore, delivers unparalleled performance on high-performance NVMe SSDs. It can achieve near-DRAM-level performance on NVMe disks available in public cloud virtual machine instances.



Core Features

10x Cost Efficiency at Scale

EloquentKV is optimized for standard Redis data structures and access patterns, but is architected around SSD-native data management. For memory-bound workloads in the **100 GB-100-TB** range per cluster, EloquentKV:

- **Offloads cold and warm data to NVMe SSDs** while keeping hot working sets in DRAM.
- **Preserves Redis-like latency characteristics**, even when the total dataset is far larger than memory.
- **Reduces infrastructure footprint by up to 10x**, often replacing large Redis clusters with a small number of NVMe-optimized nodes.

This allows teams to scale capacity based on SSD pricing, not DRAM constraints.

Primary-Database Durability and Consistency

EloquentKV is engineered as a database-grade system, not just a cache:

- **Write-Ahead Logging (WAL)** ensures durable, crash-safe persistence for every write.
- **No cache-invalidation complexity**: applications read and write to a single source of truth rather than a cache layered on top of another database.
- **Native distributed execution** supports multi-key transactions (MULTI/EXEC) and Lua scripts across shards—capabilities that are difficult or unsafe to achieve with standard Redis deployments.

- **Cluster-wide transactional semantics** maintain consistent state for cross-node transactions, enabling EloqKV to back mission-critical online systems.

Latency-Optimized and Stable Under Pressure

Conventional Redis clusters often experience latency spikes during background operations (fork-based snapshots, large RDB generation, or replica resync) and hit practical limits when per-node memory grows large. EloqKV avoids these pitfalls:

- **Zero-fork checkpointing** incrementally flushes only dirty entries, eliminating costly memory fork operations.
- **Predictable long-tail latency**: checkpointing and compaction are carefully scheduled to avoid impacting the P99-P9999 latency profile.
- **No practical per-node memory ceiling**: EloqKV comfortably supports nodes with **50 GB+** of DRAM while safely managing multi-TB datasets on SSDs.
- **Robust replica join and recovery**: standby nodes join without overwhelming primaries, avoiding the “rejoin storm” patterns of large Redis clusters (write-buffer OOM → rejoin → fail).

Fully Elastic Scaling

EloqKV is designed for full elasticity, supporting both vertical (scale-up) and horizontal (scale-out) scaling across all resources.

- **Efficient scaling up and out**: EloqKV operates efficiently at any scale. It can grow from a single node using fractional CPU cores to a fully distributed cluster with thousands of cores dynamically and without service interruption, while maintaining high performance. In addition, EloqKV preserves the same consistency and transactional semantics at all scales, eliminating the need to modify applications as the system grows.
- **Independent resource scaling**: EloqKV takes advantage of cloud platforms’ on-demand resource allocation, allowing CPU, memory, storage, and write-ahead logging (WAL) to scale independently. This enables optimal resource utilization for different workload types.

Simplified Disaster Recovery and Instant Branching

EloqKV is built for cloud-native resilience and rapid iteration:

- **Disaster Recovery (DR) via object storage**: periodic snapshots are streamed directly to cloud object storage with cross-region replication. This avoids the need for auxiliary systems like Redis

Streams or Kafka for log shipping, and keeps standby regions nearly cost-free until failover.

- **Branching in seconds:** teams can create isolated, fully functional database branches from object storage snapshots using **zero-copy cloning**. This enables realistic staging, load testing, and experimentation even on **100 TB+** datasets without duplicating physical storage.

Performance Benchmarks

EloqKV has been extensively benchmarked on NVMe SSD infrastructure to validate latency and throughput under realistic load.

- **Hardware profile:** Tests were run on Google Cloud z3-highmem-16 instances with **16 vCPUs, 128 GB RAM, and 2.9 TB NVMe SSD × 2**.
- **High throughput at DRAM and SSD scale:** Across dataset sizes from **20 GB (DRAM-sized) to 200 GB-2 TB (SSD-sized)**, EloqKV sustained **~600K QPS** for DRAM-resident access and **~300K QPS** when serving purely from disk.
- **Consistently low long-tail latency:** On a **2 TB** dataset (~800 million keys, 2.5 KB average value size) with only **100 GB of DRAM** and **100K QPS**, EloqKV maintained **P99.99 latency between 1.5 ms and 3.1 ms**, depending on the read/write mix.
- **Linear horizontal scale-out:** Throughput scales roughly linearly with node count (e.g., 10× nodes ≈ 10× throughput). EloqKV remains **compatible with Redis smart clients** and supports client-side routing across shards.

Hardware and Software Setup

The benchmark configuration is summarized below:

Service type	Node type	Node count	Disk configuration
EloqKV	z3-highmem-16	1	2.9 TB × 1 NVMe
Redis	z3-highmem-16	1	2.9 TB × 1 NVMe
KvRocks	z3-highmem-16	1	2.9 TB × 1 NVMe

Benchmark Methodology


All benchmarks used `memtier_benchmark` with the following representative commands:

```
# load data
mentier_benchmark -t 10 -c 5 -s $SERVER_PRIVATE_IP -p 6379 -n allk
# query data
mentier_benchmark -t 10 -c 5 -s $SERVER_PRIVATE_IP -p 6379 --disti
```

We ran multiple workloads to compare EloqKV against Redis and other SSD-capable systems.

Results


Benchmark 1 - EloqKV vs. Redis across dataset sizes (throughput and P99.99 latency):



Traditional DRAM-centric systems such as Redis are fundamentally limited by memory capacity. As dataset size approaches node memory, operators are forced to either overprovision DRAM or shard aggressively. In our tests, Redis delivered good performance at small sizes but could not scale beyond memory capacity, while EloqKV continued to operate smoothly by leveraging NVMe.

At **20 GB** and **100 GB**, EloqKV delivered **higher throughput than Redis** with a nearly identical P99.99 latency profile. Once the dataset exceeded memory capacity, Redis failed to handle the workload, whereas EloqKV scaled out to **2 TB** while keeping P99.99 latency stable at a few milliseconds—performance previously achievable only with DRAM-only architectures.

Benchmark 2 - EloqKV vs. Apache KvRocks (long-tail latency across workloads):



We also compared EloqKV to Apache KvRocks, another SSD-serving key-value store. Under IO-intensive workloads, KvRocks exhibited rapidly increasing P99.99 latency as background operations interfered with foreground reads (note the log scale on the chart). EloqKV, by contrast, kept P99.99 latency bounded and predictable.


To host a **2 TB** dataset with Redis, organizations typically deploy **~20 nodes** to get enough aggregate DRAM. EloqKV achieves comparable or better long-tail latency on a **single NVMe-optimized node**, reducing infrastructure cost by up to **20x** while simplifying operations substantially.

Customer Use Cases

1. AI Chat History for a Leading Smartphone Manufacturer

- **Scenario:** A top global smartphone manufacturer powers AI chat assistants across hundreds of millions of devices. Unlike traditional human-to-human chat, AI conversations generate much larger messages that bundle prompts, system context, and model outputs. Over time, this dramatically increases the volume of chat history that must be stored and retrieved in real time.

- **Challenge with previous solution:** The customer previously ran this workload on **AWS ElastiCache for Redis**. Because Redis keeps the full dataset in DRAM, the growing prompt and message history drove cluster memory requirements—and therefore cost—up sharply.
- **EloqKV deployment:** By migrating to EloqKV, the customer kept hot conversation state in memory while safely tiering the rest of the history to NVMe SSDs, all behind a Redis-compatible interface.
- **Outcome:** The customer achieved approximately **10x cost reduction** compared to their previous ElastiCache deployment, while maintaining the low-latency experience required for interactive AI chat.



4X efficiency

-30% latency

10X cost saving

2. Feature Store for Large-Scale E-Commerce Recommendations

- **Scenario:** A major e-commerce platform maintains a feature store containing **billions of product SKUs** and associated features (pricing, inventory, embeddings, and user-item interaction signals). Each recommendation request needs to fetch feature vectors for **up to 1,000 SKUs in a single call**, with a strict **P99 latency SLA of < 20 ms**.
- **Challenge with previous solution:** The feature store was previously implemented on Redis. To meet latency targets, the team stored the entire feature set in DRAM across a large Redis cluster. This resulted in very high memory costs and operational complexity as product and feature counts grew.
- **EloqKV deployment:** The company migrated the feature store to EloqKV, using Redis-compatible data structures for fast scatter-gather access patterns. Hot features remain in DRAM, while the much larger tail of SKUs is served directly from NVMe SSDs without violating latency SLAs.
- **Outcome:** EloqKV **reduced DRAM usage and cluster size by roughly 10x** while continuing to meet the **P99 < 20 ms** requirement for 1,000-SKU fetches. This enabled the team to scale the

recommendation system economically as catalog size and model complexity increased.

3. Social Graph Storage for a Large Social Media Platform

- **Scenario:** A social media company stores user friendship and follow relationships for hundreds of millions of users. The social graph currently occupies around **500 GB** and is projected to grow to **1 TB+**. Friendship data is accessed extremely frequently whenever a user is online—for feed generation, mutual-friend suggestions, and safety checks.
- **Challenges with previous solutions:**
 - With **Redis**, all online and offline friendship edges had to be kept fully in memory. This was expensive and still required a separate MySQL cluster for persistence and recovery.
 - The company later migrated to **Amazon DynamoDB**, which provided durable storage at lower day-one cost. However, as user engagement and QPS grew beyond **~30K requests per second**, the combination of read/write throughput charges and autoscaling behavior caused costs to climb steeply—eventually exceeding the projected Redis TCO at higher traffic levels.
- **EloqKV deployment:** By adopting EloqKV as the primary store for the friendship graph, the company consolidated caching and persistence into a single Redis-compatible system. Hot graph neighborhoods are served from DRAM, while the rest of the graph is efficiently stored on NVMe SSDs with database-grade durability.
- **Outcome:** The platform achieved **around 10x cost savings** compared to the projected cost of scaling either Redis or DynamoDB for this workload, while still delivering the low-latency graph lookups required for real-time social experiences at peak traffic.